

Image Classification by Convolutional Neural Networks

Aung Hein

18 Jun 2024

Outline

Introduction

Data preprocessing

VGG16

InceptionV3

ResNet152V2

Conclusion

Appendix

Introduction

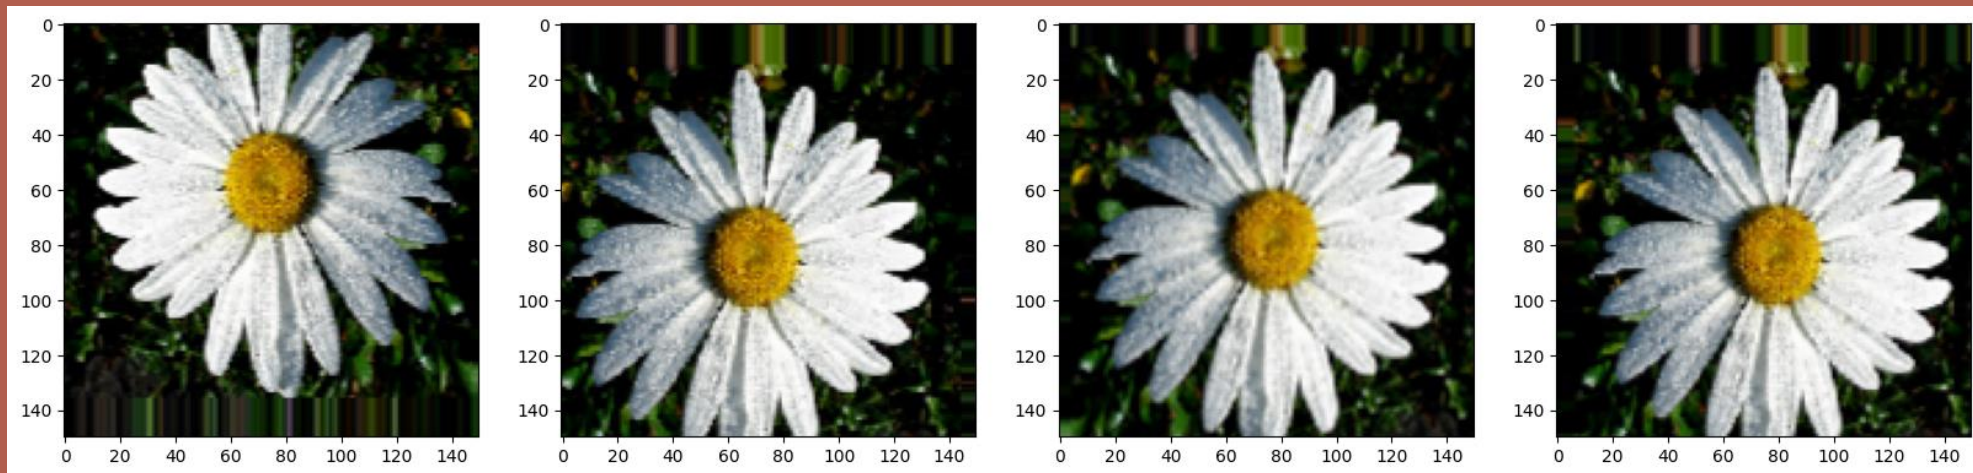
- In this project, I do experiment about few pretrained Convolutional Neural Network for image classification
- I been tested on VGG16, Inception V3 and ResNet152V2
- I used flower images dataset from TensorFlow

Dataset link:

- https://www.tensorflow.org/tutorials/load_data/images

Data Preprocessing

- There are 5 classes on image dataset : ['daisy', 'dandelion', 'rose', 'sunflower', 'tulip']
- Image data integer was rescaled by dividing with 255
- 3457 train images and 860 validation images was split, generate variations and prepared by ImageDataGenerator class





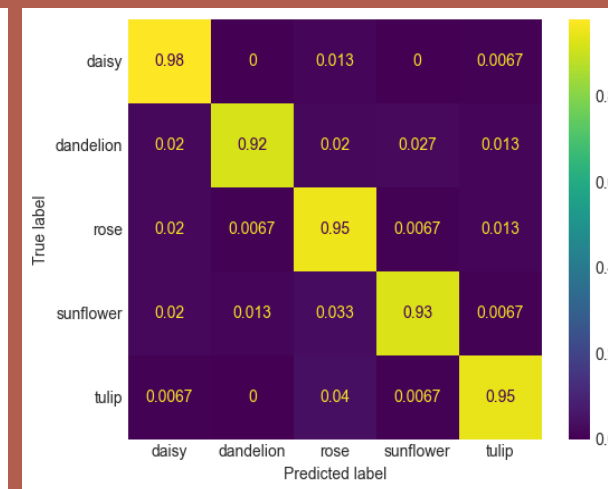


Visual Geometry Group (VGG16)

VGG16

Optimizer :

'adam'

val_accuracy:
0.7686




Model: "sequential"

Layer (type)	Output Shape	Param #
model (Functional)	(None, 8192)	14714688
dense (Dense)	(None, 5)	40965

=====
Total params: 14,755,653
Trainable params: 40,965
Non-trainable params: 14,714,688

Image with Predicted Labels






VGG16 with additional dense layer

Optimizer :

'adam'

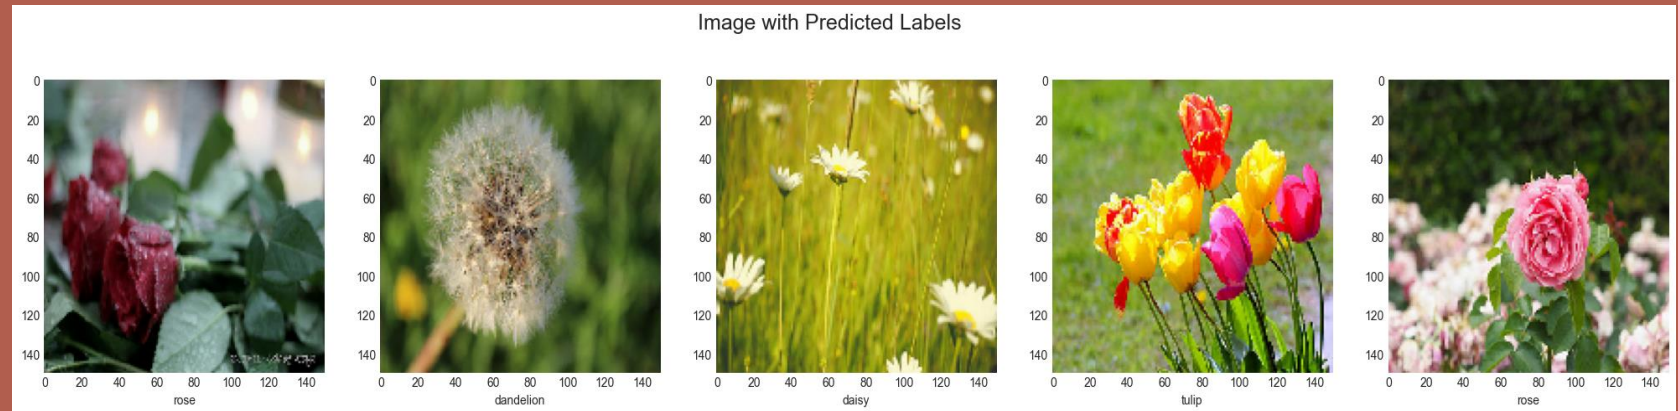
val_accuracy:
0.7791



Model: "sequential_6"

Layer (type)	Output Shape	Param #
model_2 (Functional)	(None, 8192)	14714688
dense_18 (Dense)	(None, 1048)	8586264
dropout_6 (Dropout)	(None, 1048)	0
dense_19 (Dense)	(None, 128)	134272
dense_20 (Dense)	(None, 5)	645

Total params: 23,435,869
Trainable params: 8,721,181
Non-trainable params: 14,714,688



VGG16 with Tuning

Optimizer :

SGD (lr = 0.001)

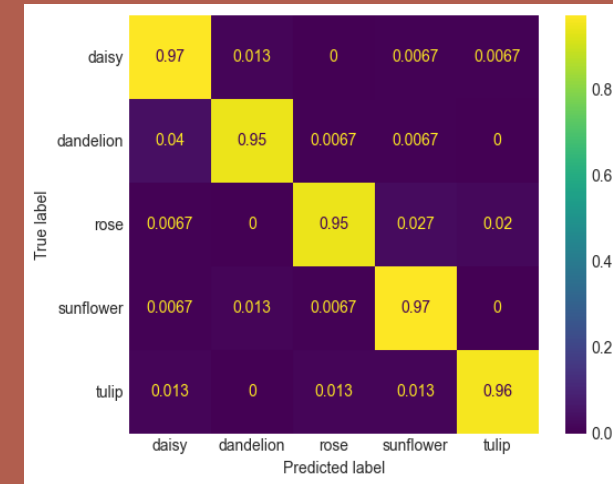


val_accuracy:
0.8628

Tuned layer :

'block5_conv1'

'block5_conv2'

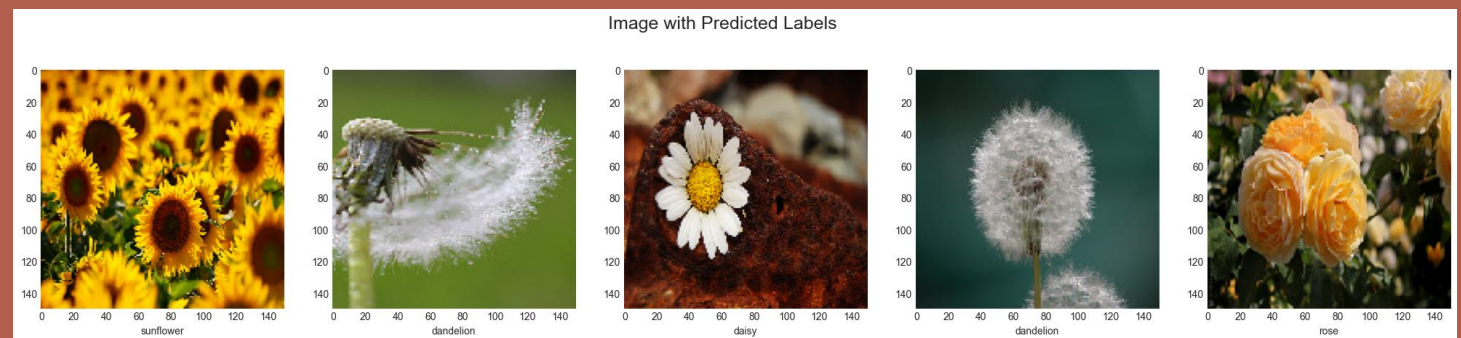
'block5_conv3'

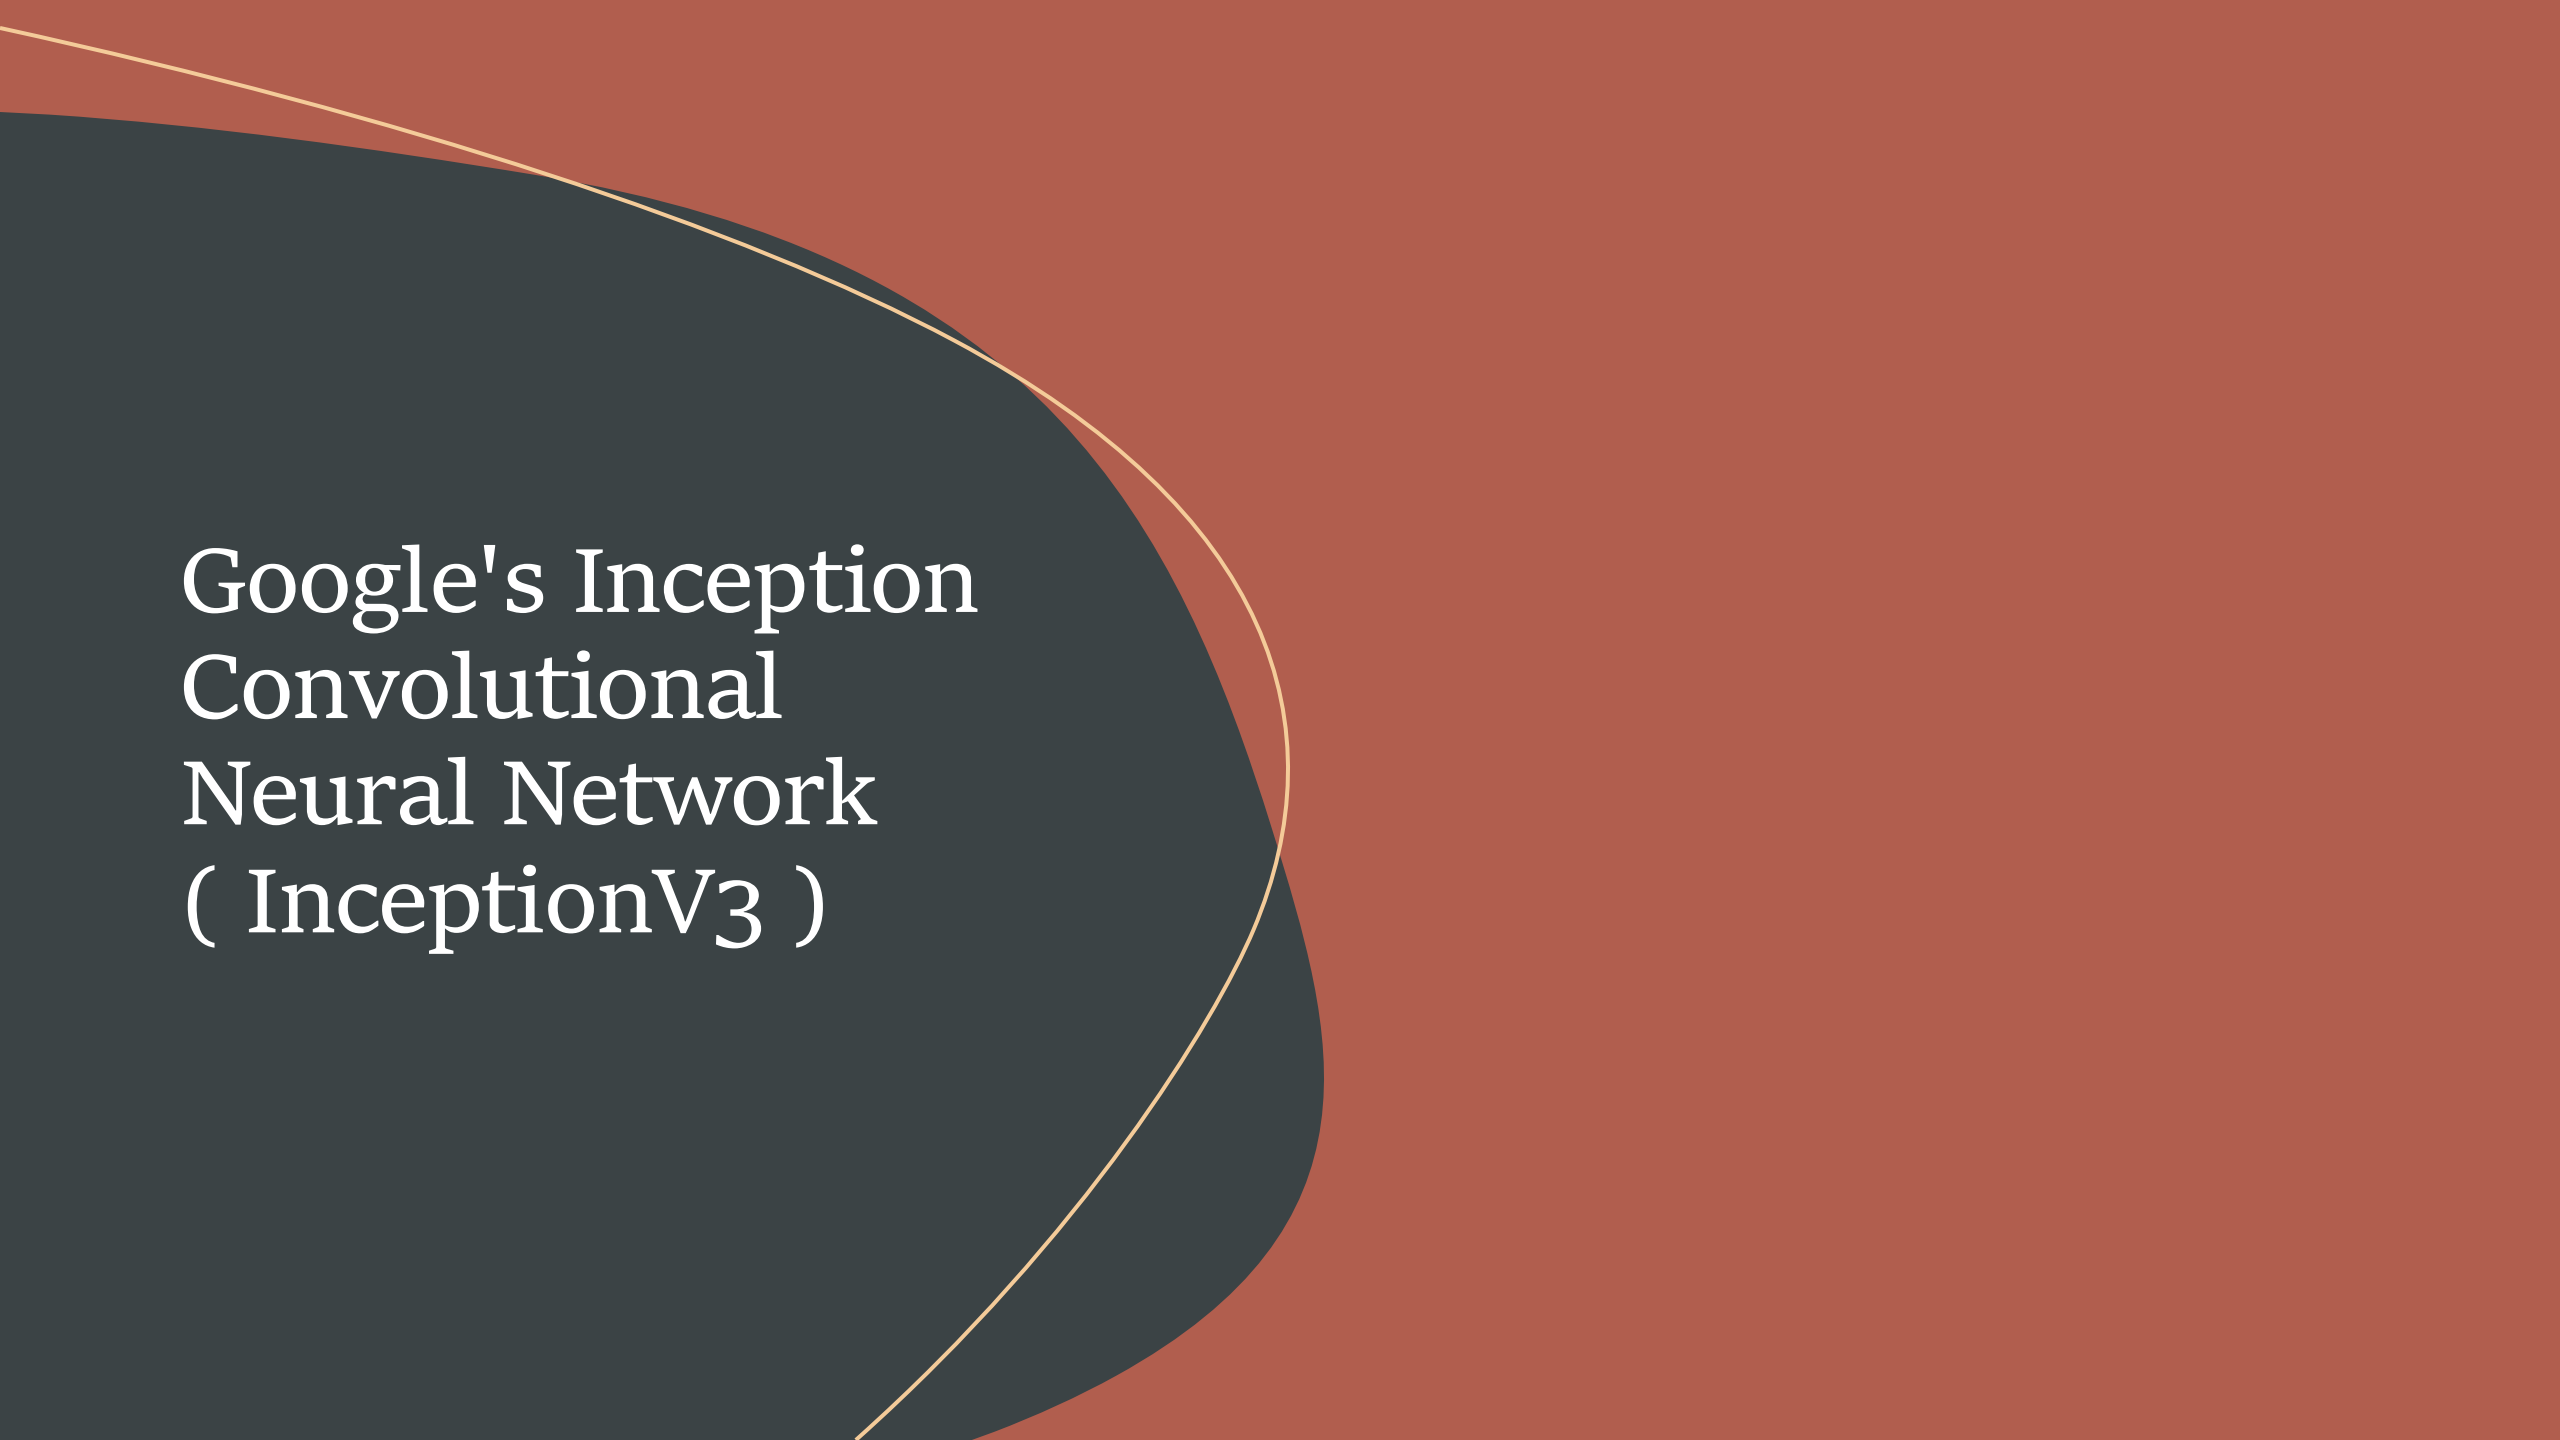


Model: "sequential_7"

Layer (type)	Output Shape	Param #
model_2 (Functional)	(None, 8192)	14714688
dense_21 (Dense)	(None, 1048)	8586264
dropout_7 (Dropout)	(None, 1048)	0
dense_22 (Dense)	(None, 128)	134272
dense_23 (Dense)	(None, 5)	645

Total params: 23,435,869
Trainable params: 15,800,605
Non-trainable params: 7,635,264





Google's Inception Convolutional Neural Network (InceptionV3)

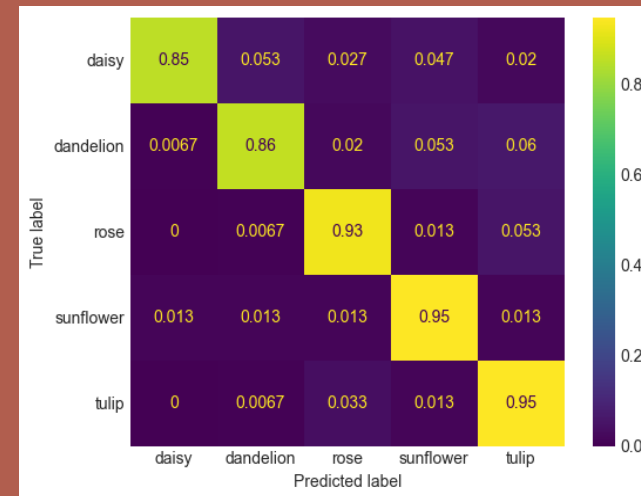


InceptionV3

Optimizer :

Adam

val_accuracy:

0.7919



Model: "sequential"

Layer (type)	Output Shape	Param #
model (Functional)	(None, 18432)	21802784
dense (Dense)	(None, 5)	92165

Total params: 21,894,949

Trainable params: 92,165

Non-trainable params: 21,802,784



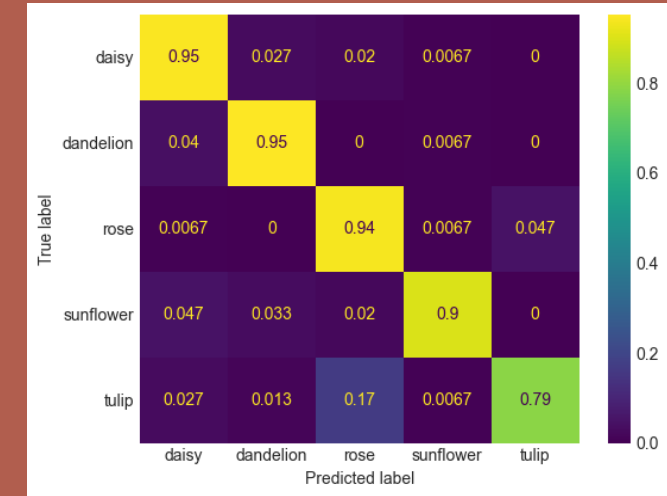


InceptionV3 with additional dense layer

Optimizer :

Adam

val_accuracy:

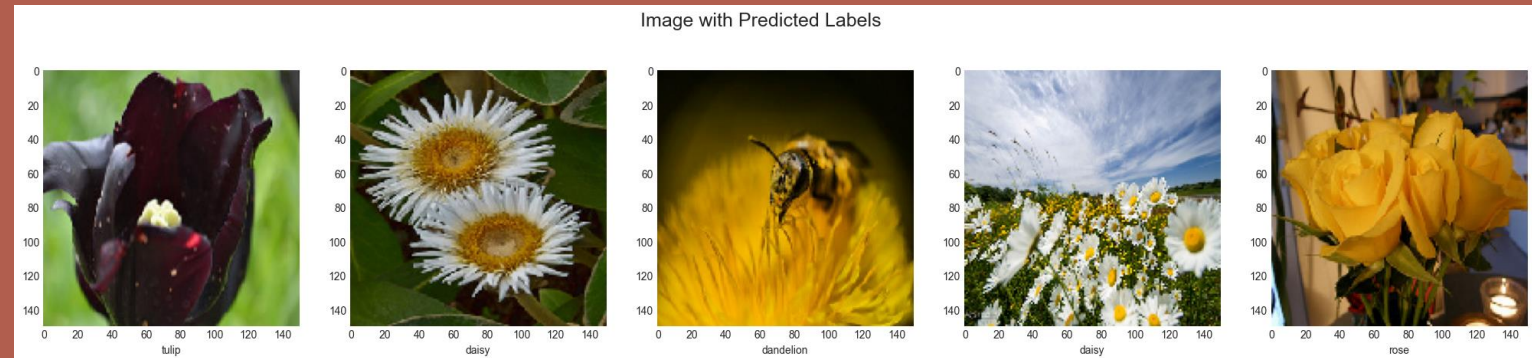
0.7953



Model: "sequential"

Layer (type)	Output Shape	Param #
model (Functional)	(None, 18432)	21802784
dense (Dense)	(None, 1048)	19317784
dropout (Dropout)	(None, 1048)	0
dense_1 (Dense)	(None, 128)	134272
dense_2 (Dense)	(None, 5)	645

Total params: 41,255,485
Trainable params: 19,452,701
Non-trainable params: 21,802,784



InceptionV3 with Tuning

Optimizer :

SGD (lr = 0.001)




val_accuracy:
0.8372

Tuned layer :

'conv2d_89',

'conv2d_86',

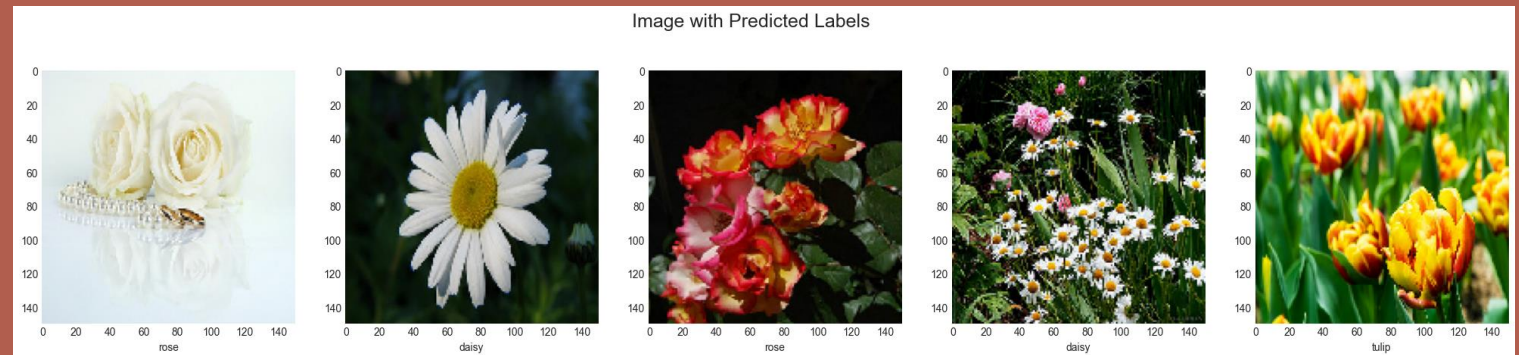
'conv2d_85'

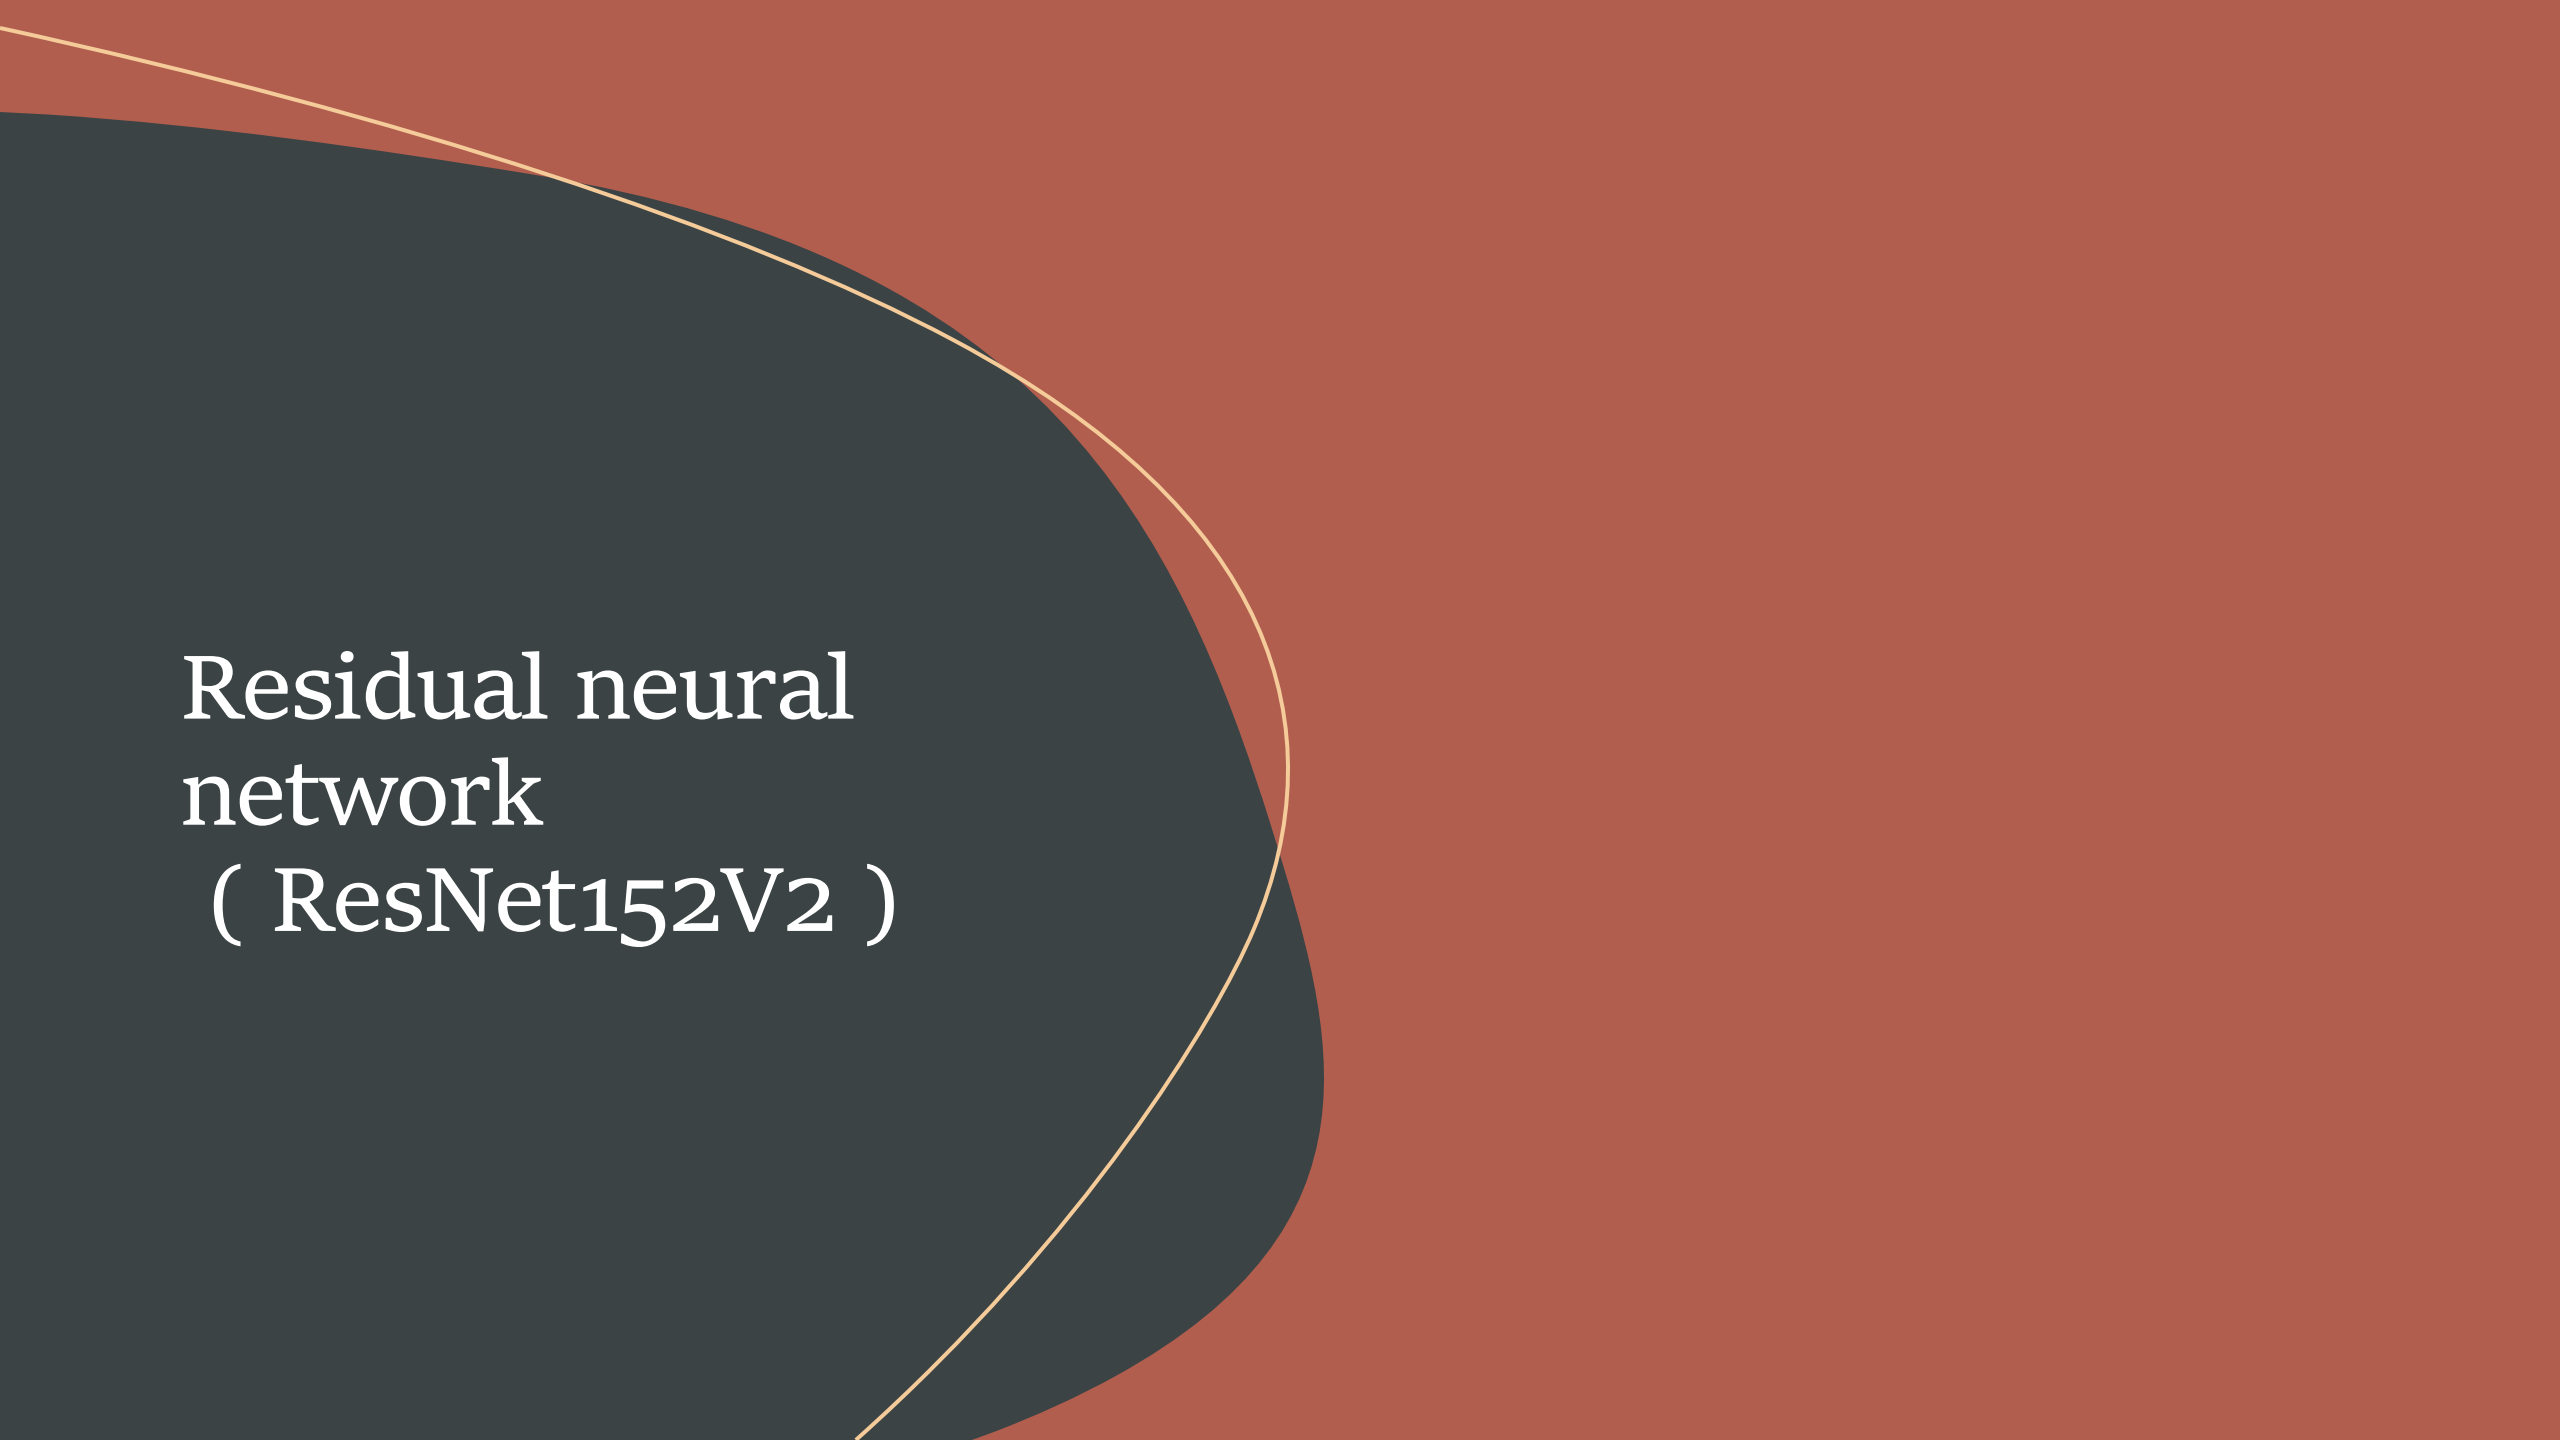


Model: "sequential"

Layer (type)	Output Shape	Param #
model (Functional)	(None, 18432)	21802784
dense (Dense)	(None, 1048)	19317784
dropout (Dropout)	(None, 1048)	0
dense_1 (Dense)	(None, 128)	134272
dense_2 (Dense)	(None, 5)	645

Total params: 41,255,485
Trainable params: 25,526,237
Non-trainable params: 15,729,248





Residual neural network (ResNet152V2)




ResNet152V2

Optimizer :

Adam

val_accuracy:

0.7919



Model: "sequential"

Layer (type)	Output Shape	Param #
model (Functional)	(None, 18432)	21802784
dense (Dense)	(None, 5)	92165

Total params: 21,894,949

Trainable params: 92,165

Non-trainable params: 21,802,784



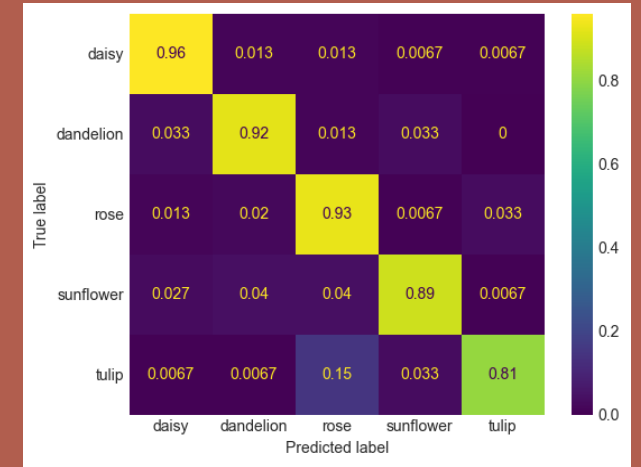


ResNet152V2 with additional dense layer

Optimizer :

Adam

val_accuracy:

0.8233




Model: "sequential"

Layer (type)	Output Shape	Param #
model (Functional)	(None, 51200)	58331648
dense (Dense)	(None, 1048)	53658648
dropout (Dropout)	(None, 1048)	0
dense_1 (Dense)	(None, 128)	134272
dense_2 (Dense)	(None, 5)	645

Total params: 112,125,213
Trainable params: 53,793,565
Non-trainable params: 58,331,648

Image with Predicted Labels



ResNet152V2 with Tuning




Optimizer :

SGD (lr = 0.001)

val_accuracy:
0.8419

Tuned layer :

'conv5_block3_3_
conv',
'conv5_block3_2_
conv'



Model: "sequential"

Layer (type)	Output Shape	Param #
model (Functional)	(None, 51200)	58331648
dense (Dense)	(None, 1048)	53658648
dropout (Dropout)	(None, 1048)	0
dense_1 (Dense)	(None, 128)	134272
dense_2 (Dense)	(None, 5)	645



Total params: 112,125,213

Trainable params: 57,208,605

Non-trainable params: 54,916,608



Conclusion



- By adding more hidden layers or complexity, it always got slighter improve result.
- By tunning, later few layers, we get a huge improvement and got the best result.

Appendix

Github Link : <https://github.com/thonenyangal/Image-Classification.git>